

Efficient R

Selina Baldauf

✉ selina.baldauf@fu-berlin.de 🦋 @selina-b

Welcome

- Ecology/Theoretical Ecology
- Scientific programmer at Freie Universität Berlin
- Research + Coding + Teaching



What is efficiency?

$$\text{efficiency} = \frac{\text{work done}}{\text{unit of effort}}$$

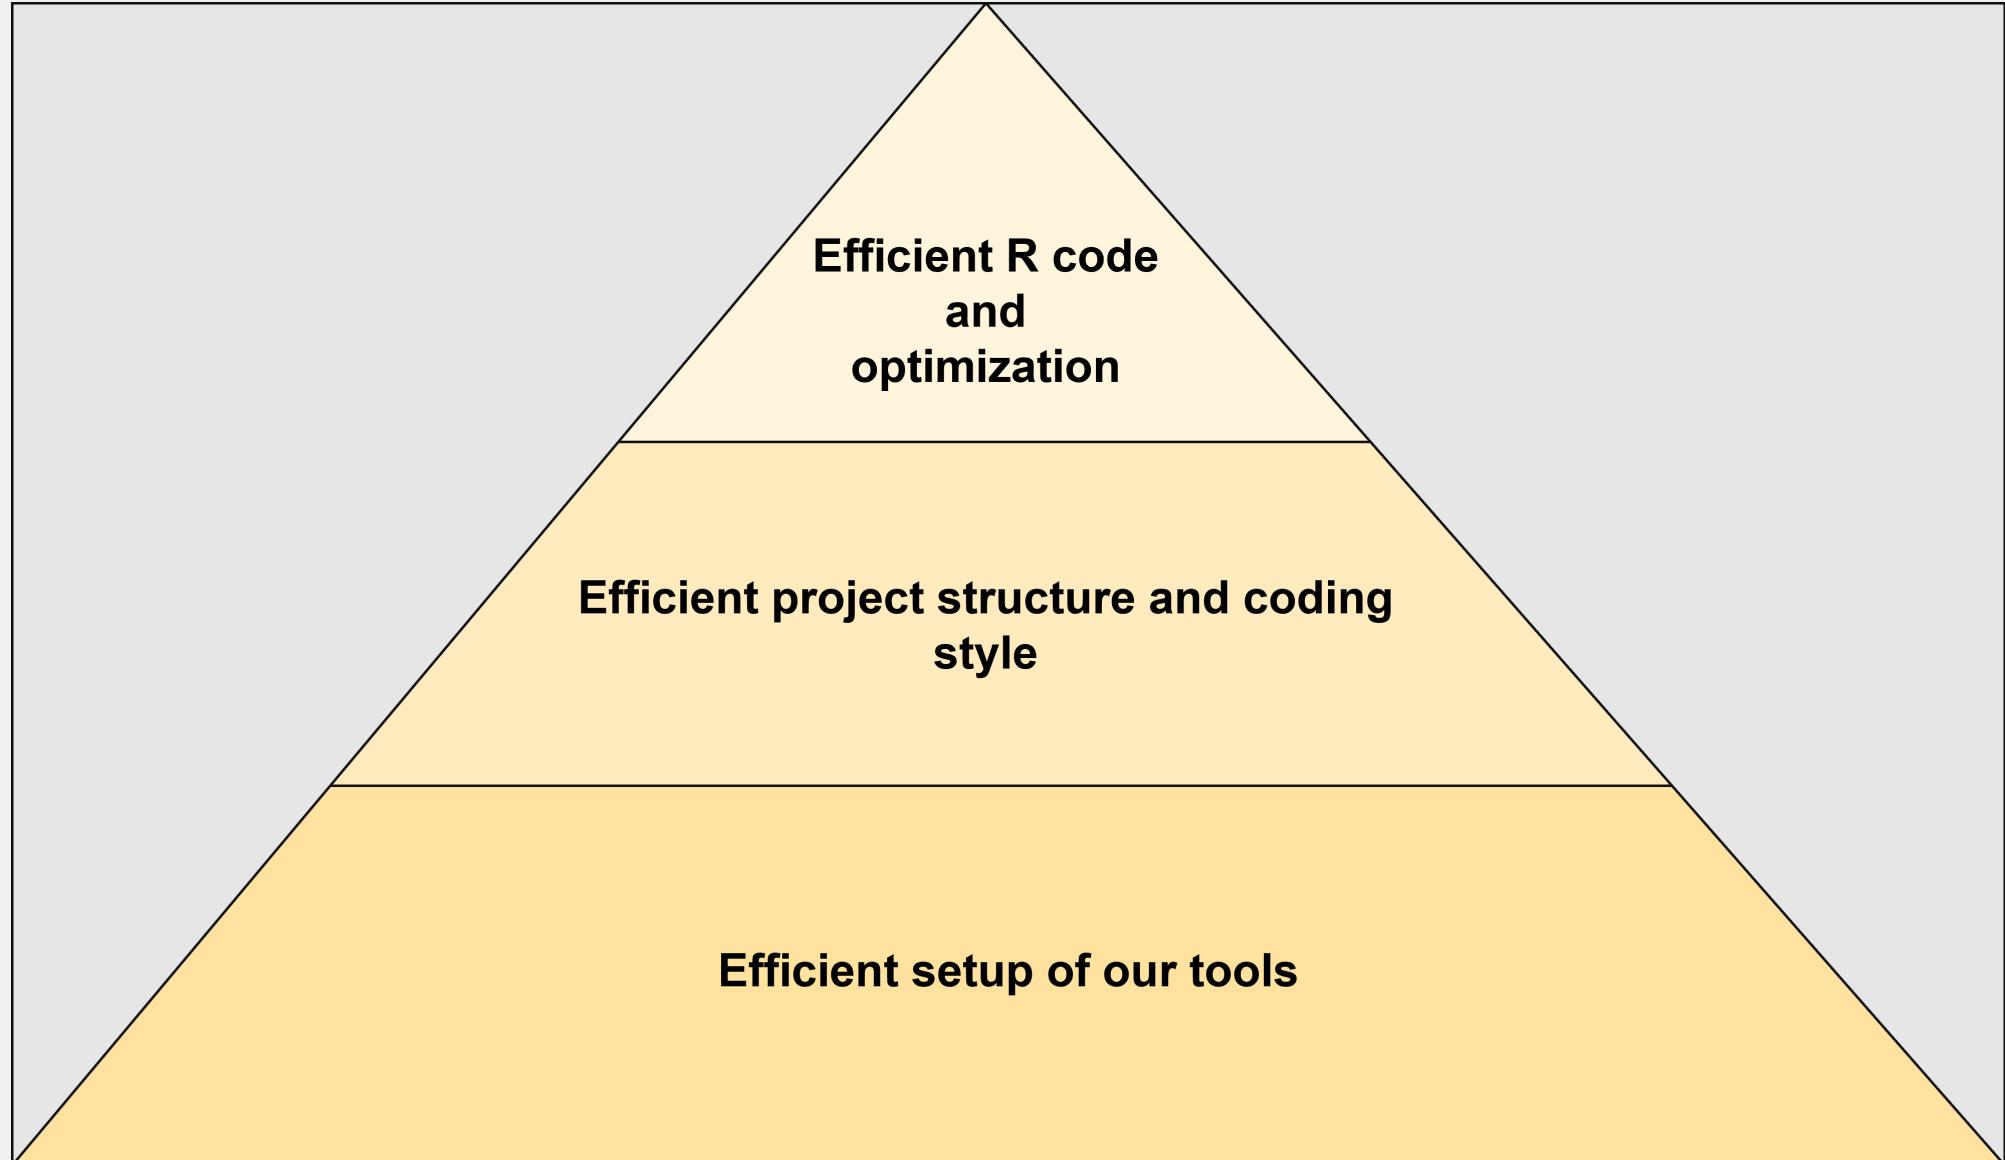
Computational efficiency

-  Computation time
-  Memory usage

Programmer efficiency

-  How long does it take to
 - *write* code?
 - *Maintain* code?
 - *read* and *understand* the code?

Tradeoffs and Synergies between these types of efficiencies



Principles and tools to make R programming more efficient for the 

Today

- How to **measure speed** of your code
 - **Basics** of efficient R programming
 - Efficient **data analysis**
 - Advanced optimization
 - Parallelization
 - Integrating C++
-  Material (Slides, Code, References & Resources) on [Github](#)
-  Question - just ask in the Chat or unmute yourself!

Is R slow?

- R is slow compared to other programming languages (e.g. C++, Julia).
- R is not the most memory efficient language
- R is designed to make statistical programming & data analysis **easy** and **interactive**, not fast
- But: **R is fast and memory efficient enough** for most tasks.

Should I optimize?

It's easy to get caught up in trying to remove all bottlenecks. Don't! Your time is valuable and is better spent analysing your data, not eliminating possible inefficiencies in your code. Be pragmatic: don't spend hours of your time to save seconds of computer time.
(Hadley Wickham in *Advanced R*)

Think about

- How much time do I **save** vs. **spend** optimizing?
- **How often** do I run the code?
- Trade-offs between **readability** and **efficiency**

How to optimize

Don't optimize from top to bottom!

1. Identify bottlenecks (i.e. slow parts) of your code
2. Optimize only those bottlenecks

Profiling and benchmarking

Measure the speed and memory usage of your code

Profiling R code

What are the speed & memory **bottlenecks** in my code?

Use the `profvis` package

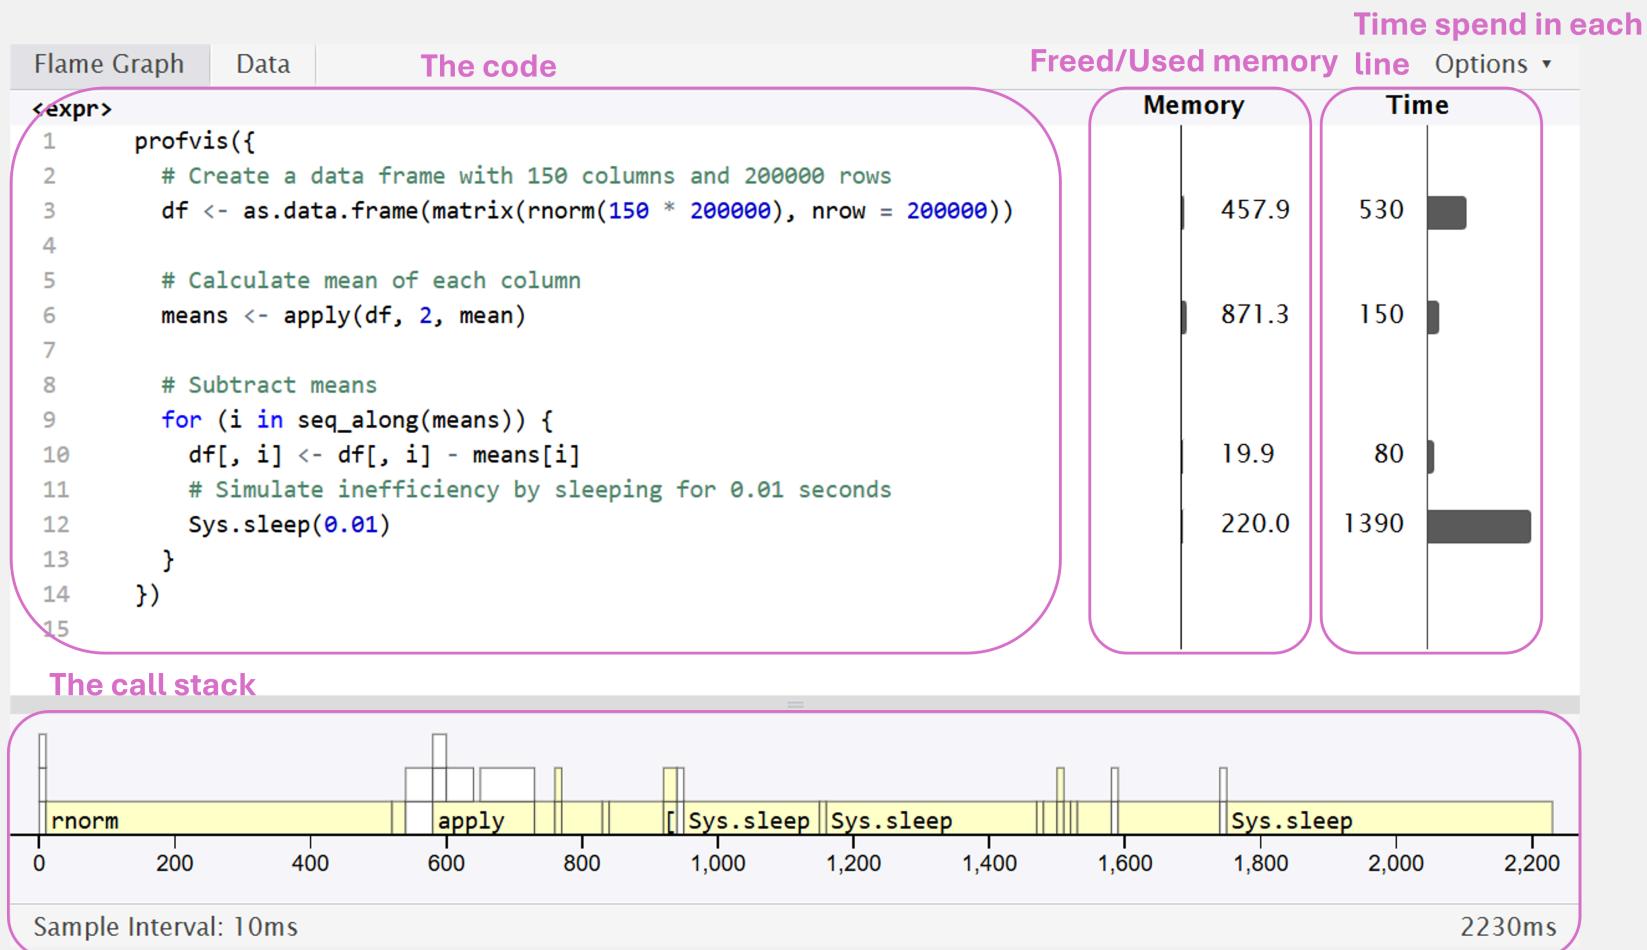
Profiling R code

You can profile a section of code like this:

```
1 # install.packages("profvis")
2 library(profvis)
3
4 profvis({
5   # Create a data frame with 150 columns and 200000 rows
6   df <- as.data.frame(matrix(rnorm(150 * 200000), nrow = 200000))
7
8   # Calculate mean of each column
9   means <- apply(df, 2, mean)
10
11  # Subtract means
12  for (i in seq_along(means)) {
13    df[, i] <- df[, i] - means[i]
14    # Simulate inefficiency by sleeping for 0.01 seconds
15    Sys.sleep(0.01)
16  }
17 })
```

Profiling R code

Profvis flame graph shows time and memory spent in each line of code.



Profiling R code

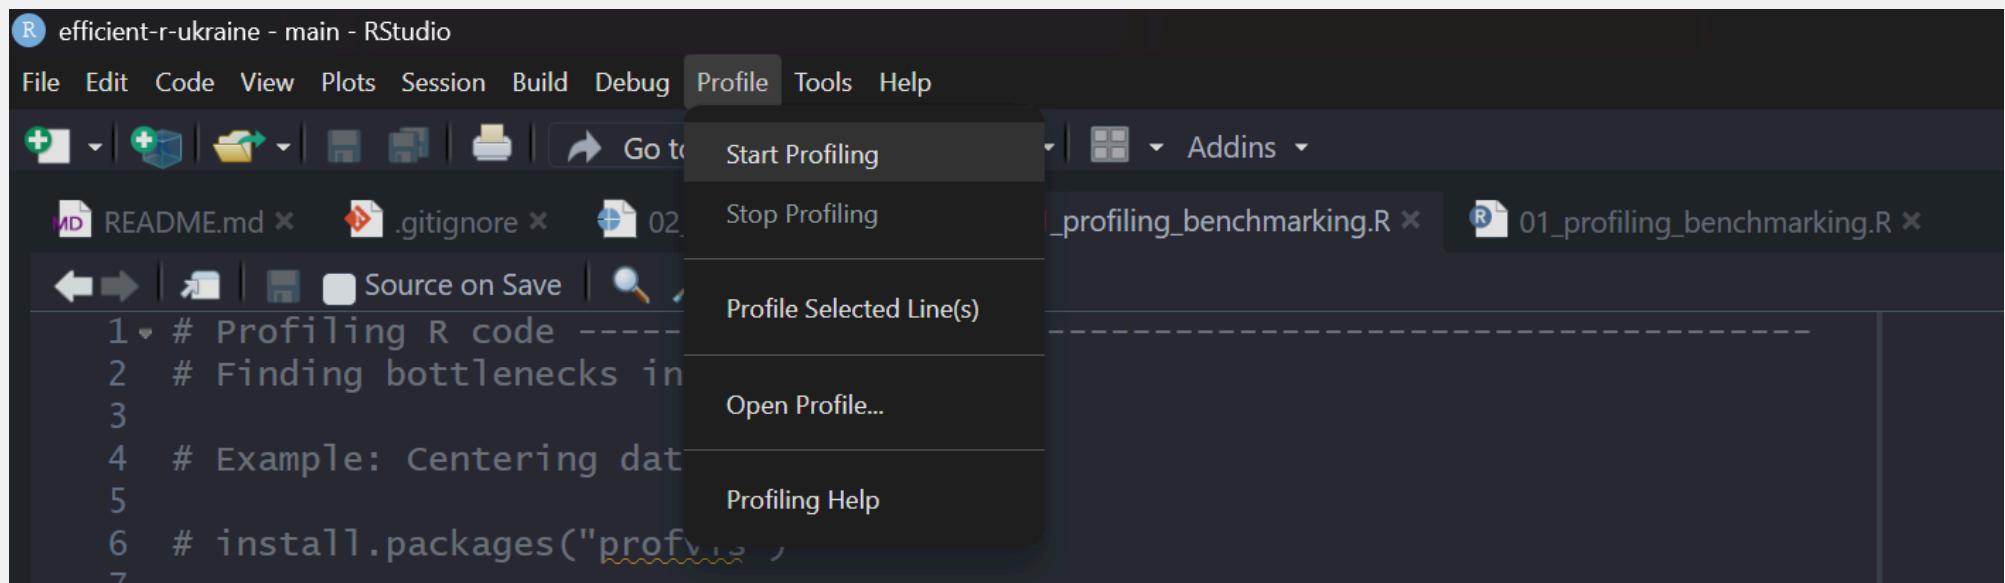
Profvis **data view** for details of the call stack.

Code	File	Memory (MB)		Time (ms)
Sys.sleep	<expr>	0	220.0	1390
rnorm	<expr>	0	229.0	510
▼ apply	<expr>	0	871.3	150
mean.default		0	184.6	80
aperm.default		0	228.9	40
as.matrix.data.frame		0	308.2	20
► profvis		0	4.6	40
as.data.frame.matrix		0	264.0	40
[<expr>	0	6.1	40
matrix	<expr>	0	228.9	20
[.data.frame	<expr>	0	3.1	20
df[, i] <- df[, i] - means[i]	<expr>	0	4.6	20

Profiling R code

You can also interactively profile code in RStudio:

- Go to **Profile -> Start profiling**
- Now interactively run the code you want to profile
- Go to **Profile -> Stop profiling** to see the results



Benchmarking R code

Which version of the code is faster?

```
center_data_slow <- function() {  
  # Create data with 150 columns and 100000 rows  
  df <- as.data.frame(  
    matrix(rnorm(150 * 100000), nrow = 100000)  
  )  
  
  # Calculate mean of each column  
  means <- apply(df, 2, mean)  
  
  # Subtract means  
  for (i in seq_along(means)) {  
    df[, i] <- df[, i] - means[i]  
    # Simulate inefficiency by sleeping 0.01 seconds  
    Sys.sleep(0.01)  
  }  
  return(df)  
}
```

```
center_data_fast <- function() {  
  # Create data with 150 columns and 100000 rows  
  df <- matrix(rnorm(150 * 100000), nrow = 100000)  
  
  # Calculate mean of each column  
  means <- colMeans(df)  
  
  # Subtract means  
  for (i in seq_along(means)) {  
    df[, i] <- df[, i] - means[i]  
  }  
  return(df)  
}
```

Benchmarking R code - the easy way

Use the `tictoc` package also for longer code sections:

```
# install.packages("tictoc")
library(tictoc)

tic()
slow_data <- center_data_slow()
toc()
#> 3.06 sec elapsed

tic()
fast_data <- center_data_fast()
toc()
#> 0.55 sec elapsed
```

Benchmarking R code

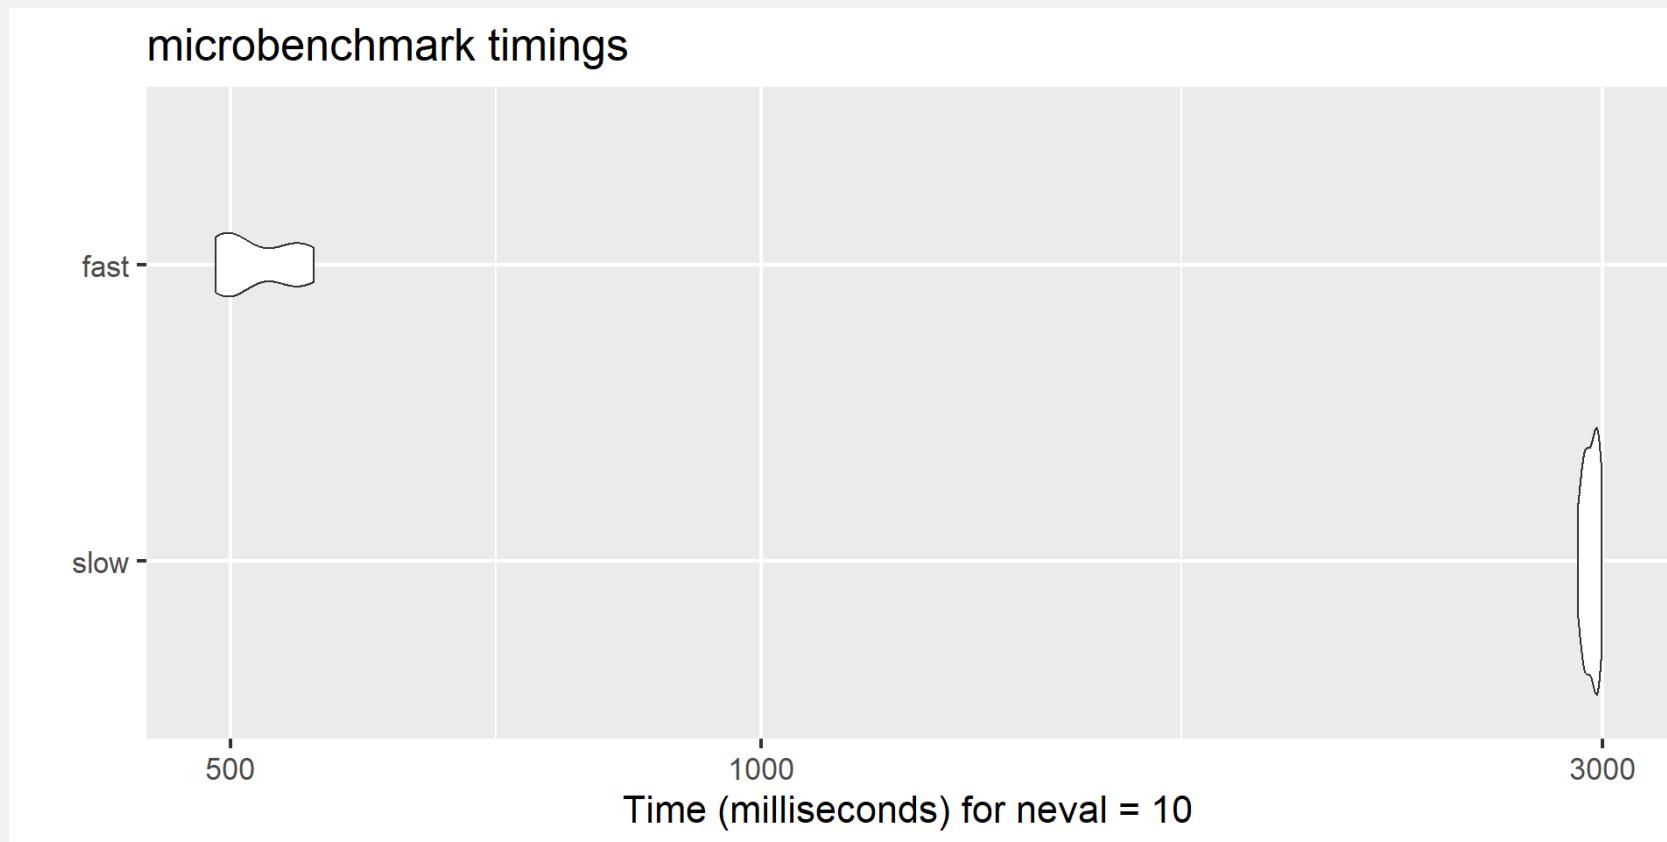
Use the `microbenchmark` package to compare the functions:

```
1 # install.packages("microbenchmark")
2 library(microbenchmark)
3
4 runtime_comp <- microbenchmark(
5   slow = center_data_slow(),
6   fast = center_data_fast(),
7   times = 10 # the default is 100 but we are impatient
8 )
9
10 runtime_comp
11 #> Unit: milliseconds
12 #>   expr      min       lq     mean    median       uq      max neval cld
13 #>   slow 2907.6015 2930.4255 2956.9953 2960.0167 2981.8489 2997.6368    10    a
14 #>   fast  490.1377  496.1393  517.5706  505.7072  542.5332  557.1841    10    b
15
16 # look at relative runtime comparison
17 summary(runtime_comp, unit = "relative")
18 #>   expr      min       lq     mean    median       uq      max neval cld
19 #> 1 slow 5.932214 5.906457 5.713222 5.853222 5.496159 5.379975    10    a
20 #> 2 fast 1.000000 1.000000 1.000000 1.000000 1.000000 1.000000    10    b
```

Benchmarking R code

We can look at benchmarking results using ggplot

```
library(ggplot2)
autoplot(runtime_comp)
```



Optimize your R code

- Basic principles
- Data analysis bottlenecks
- Advanced optimization: Parallelization and C++

Basic principles

Vectorize your code

- Vectors are central to R programming
- R is optimized for vectorized code
 - Implemented directly in C/Fortran
- Vector operations can often replace for-loops in R
- If there is a vectorized version of a function: Use it

Vectorized functions in R

Some examples of base-R vectorized functions:

- Arithmetic & Math: `+`, `-`, `*`, `/`, `log()`, `exp()`, `sqrt()`, `mean()`, `sd()`, `rowSums()`, `colMeans()`
- Logical & Comparison: `<`, `>`, `==`, `!=`, `&`, `|`, `all()`, `any()`, `%in%`
- Data Manipulation: `[`, `ifelse()`, `cut()`, `cumsum()`, `which()`, `match()`, `order()`
- String operations: `paste()`, `paste0()`, `grep()`, `gsub()`

Vectorize your code

Vector arithmetic vs. for-loops

Example 2: Calculating cumulative values

```
x <- 1:1e6
y <- 1:1e6

microbenchmark(
  for_loop = {
    result <- numeric()
    for (i in seq_along(x)) {
      result[i] <- x[i] * 2 + y[i] / 2
    }
  },
  vectorized = x * 2 + y / 2,
  times = 10
)
#> Unit: milliseconds
#>          expr      min       lq     mean   median      uq     max neval cld
#>  for_loop 121.2896 125.0952 139.65615 129.18735 136.9957 190.5542     10    a
#>  vectorized  1.8256  1.8668  2.89729  3.22085  3.6280  3.6639     10    b
```

For-loops in R

- For-loops are **relatively slow** and it is easy to make them even slower with bad design
- Often they are used when vectorized code would be better
- For loops can often be replaced, e.g. by
 - Functions from the `apply` family (e.g. `apply`, `lapply`, ...)
 - Vectorized functions (e.g. `sum`, `colMeans`, ...)
 - Vectorized functions from the `purrr` package (e.g. `map`)

But: For loops are not necessarily bad, **sometimes** they are the **best solution** and **more readable** than vectorized code.

Don't grow objects in a loop

If you know how big your object will be, pre-allocate it.

```
f1 <- function() {  
  x <- numeric() # no pre-allocation  
  for (i in 1:1e6) {  
    x[i] <- i  
  }  
}
```

```
f2 <- function() {  
  x <- numeric(1e6) # pre-allocate vector  
  for (i in 1:1e6) {  
    x[i] <- i  
  }  
}
```

```
compare_alloc <- microbenchmark(  
  no_alloc = f1(),  
  pre_alloc = f2(),  
  times = 10  
)  
  
summary(compare_alloc, unit = "relative")  
#>      expr      min       lq     mean   median      uq     max neval cld  
#> 1  no_alloc 1.0000000 1.0000000 1.0000000 1.000000 1.0000000    10    a  
#> 2 pre_alloc 0.9650891 0.9896431 0.9831439 1.009073 1.017852 0.8558446    10    a
```

Don't grow objects in a loop

Why?

R makes a copy of the object each time it grows, which is inefficient.

Don't grow objects in a loop

```
x <- numeric()
y <- numeric(10)

for (i in 1:10) {
  x[i] <- i
  y[i] <- i

  # Print current memory location after each assignment
  cat(
    "After iteration ",
    i,
    ": x is in ",
    pryr::address(x),
    "and y is in ",
    pryr::address(y),
    "\n"
  )
}

#> After iteration 1 : x is in 0x1c07a517498 and y is in 0x1c06627cc38
#> After iteration 2 : x is in 0x1c065e5f5c8 and y is in 0x1c06627cc38
#> After iteration 3 : x is in 0x1c0928940a8 and y is in 0x1c06627cc38
```

Cache variables

If you use a value multiple times, store it in a variable to avoid re-calculation

Example: Calculate column means and normalize them by the standard deviation

```
1 # A matrix with 1000 columns
2 x <- matrix(rnorm(10000), ncol = 1000)
3
4 microbenchmark(
5   no_cache = apply(x, 2, function(i) mean(i) / sd(x)),
6   cache = {
7     sd_x <- sd(x)
8     apply(x, 2, function(i) mean(i) / sd_x)
9   }
10 )
11 #> Unit: milliseconds
12 #>      expr      min       lq     mean   median      uq      max neval cld
13 #> no_cache 52.1793 56.61065 68.90352 72.38950 75.22725 94.6167    100   a
14 #>     cache  3.3721  3.58560  3.97894  3.65755  3.77115 19.7137    100   b
```

Efficient data analysis

Efficient workflow

- Filter early: Remove unnecessary rows/columns at the beginning
- Avoid redundant calculations
 - Calculate values once and store them in variables
 - Save intermediate results
- Use efficient data formats
 - Text files instead of Excel files
 - Binary formats like Parquet and feather for large datasets
 - Consider databases for very large datasets
- Use the right packages and functions for your tasks

tidyverse

- Collection of packages for data analysis
- Consistent syntax, works well with pipes (`|>`, `%>%`)
- Main data structure `tbl_df`
- Universe of efficient packages build around the core tidyverse, e.g.
 - `dbplyr` for database and `dtplyr` for data.table backends
 - `futureverse` packages for parallelization
- Often fast enough but can be slower than alternatives

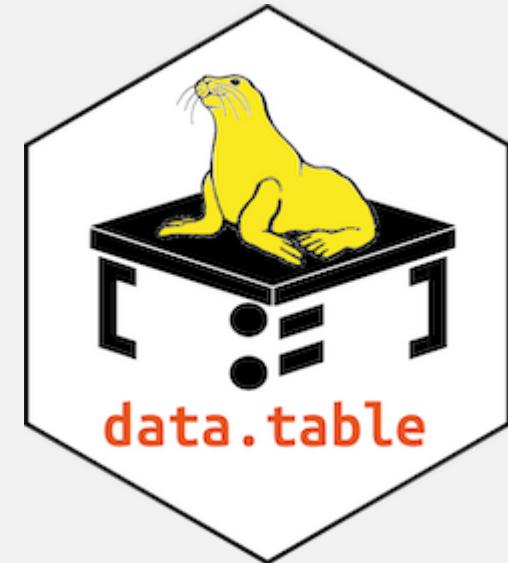


data.table

- Functions for I/O, data manipulation and more
- Popular alternative to the `tidyverse`
- Designed for fast and memory-efficient data analysis
- Syntax quite different from tidyverse:

```
dt[i, j, by]
```

Take `data.table` `dt`, subset rows using `i` and manipulate columns with `j`, grouped according to `by`.



collapse

collapse is written in C and C++, with algorithms much faster than base R's, has extremely low evaluation overheads, scales well, and excels on complex statistical tasks.

- Data manipulation and statistical computing
- Easy to integrate with existing R code
 - works well with both `data.table` and `dplyr` workflows/data structures
 - often prefixes functions with `f` like `fmean`, `fmutate`, `fsummarize`, ...



Arrow

- Access to features of Apache Arrow C++ library
- Functions for I/O and data manipulation
- Access to efficient binary data formats like parquet, feather, etc.
- Can also operate on larger-than-RAM data
- Provide Arrow C++ backend to `dplyr`



Other packages

Of course there are so many more packages like:

- `polars`: Fast DataFrame library implemented in Rust
- `fst`: Fast and efficient data format with high read/write speed
- `memoise`: Cache function calls
- ...

Do you know others? Let's collect them in the chat

Check the resources section of the README for all links.

Read data - text and excel

Example: Read data on global greenhouse gas emissions (~320000 rows, 8 cols).

```
file_path_csv <- here::here("data/ghg_emissions.csv")
file_path_xlsx <- here::here("data/ghg_emissions.xlsx")

compare_read <- microbenchmark(
  # base R
  read.csv = read.csv(file_path_csv),

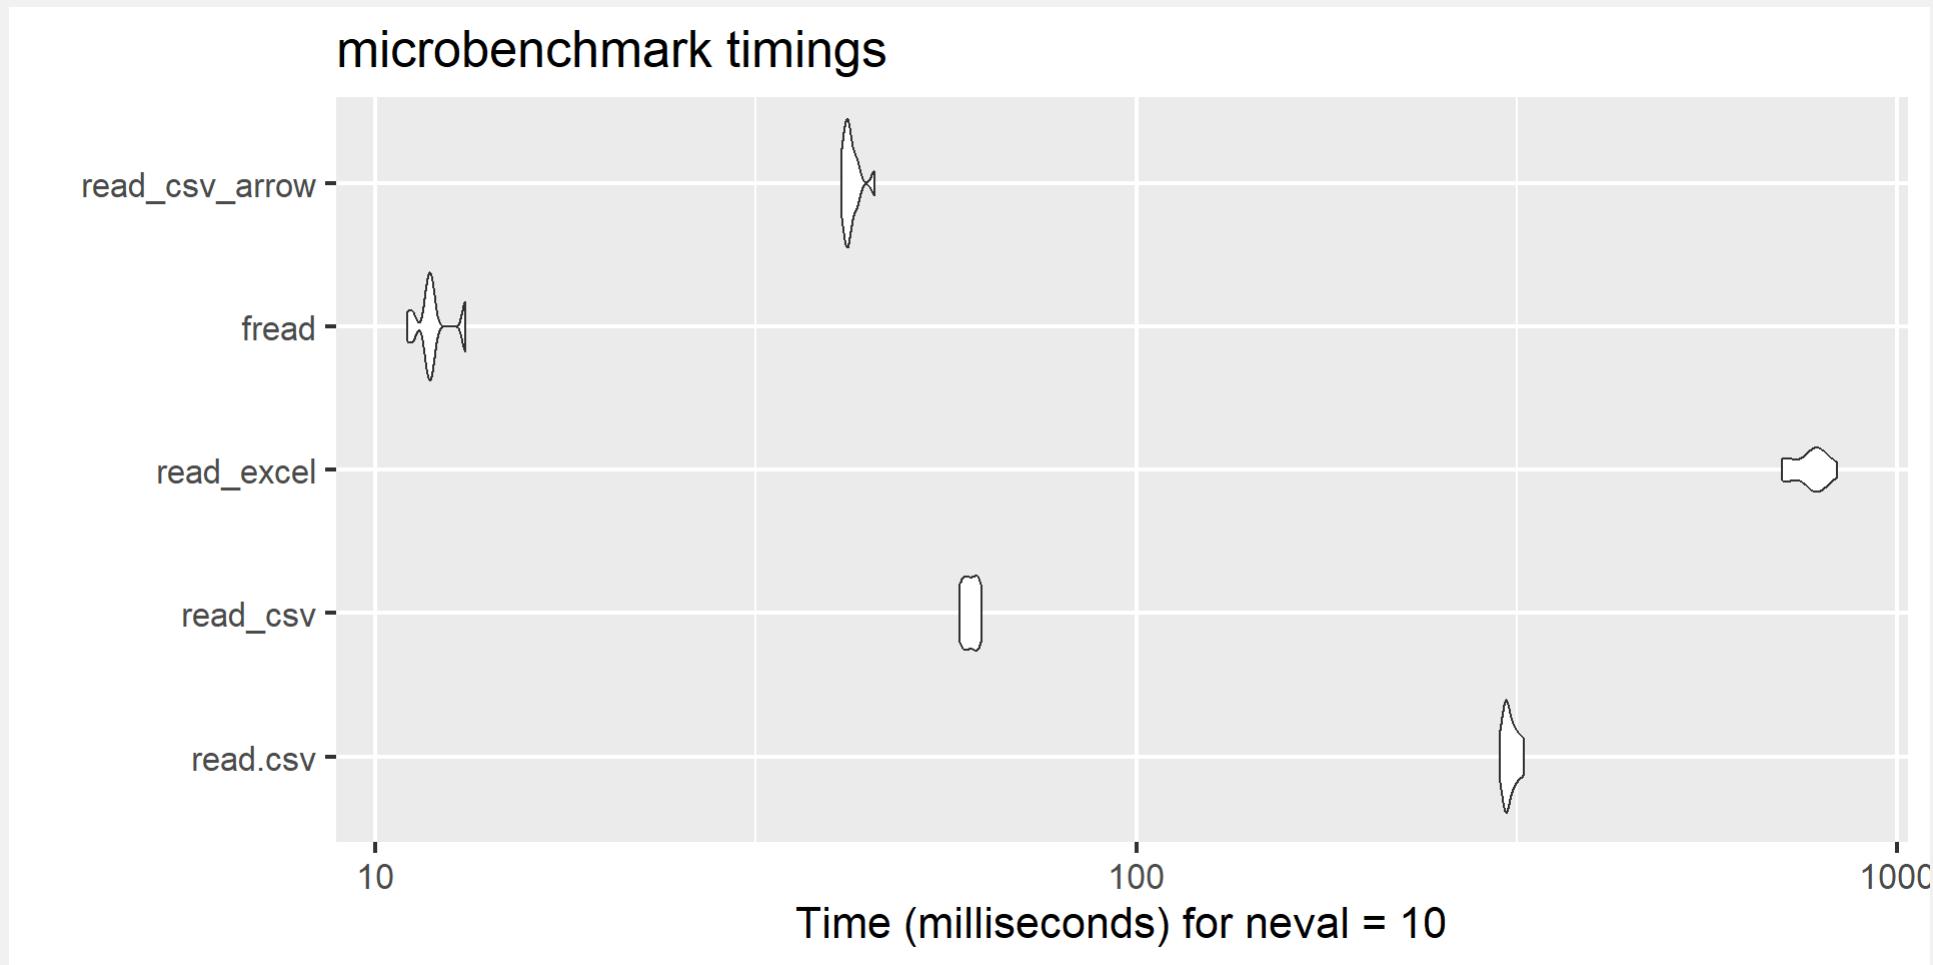
  # tidyverse
  read_csv = readr::read_csv(file_path_csv, show_col_types = FALSE),
  read_excel = readxl::read_excel(file_path_xlsx),

  # data.table
  fread = data.table::fread(file_path_csv, showProgress = FALSE),

  # arrow
  read_csv_arrow = arrow::read_csv_arrow(file_path_csv),
  times = 10
)

autoplot(compare_read)
```

Read data - text and excel



Read data - binary

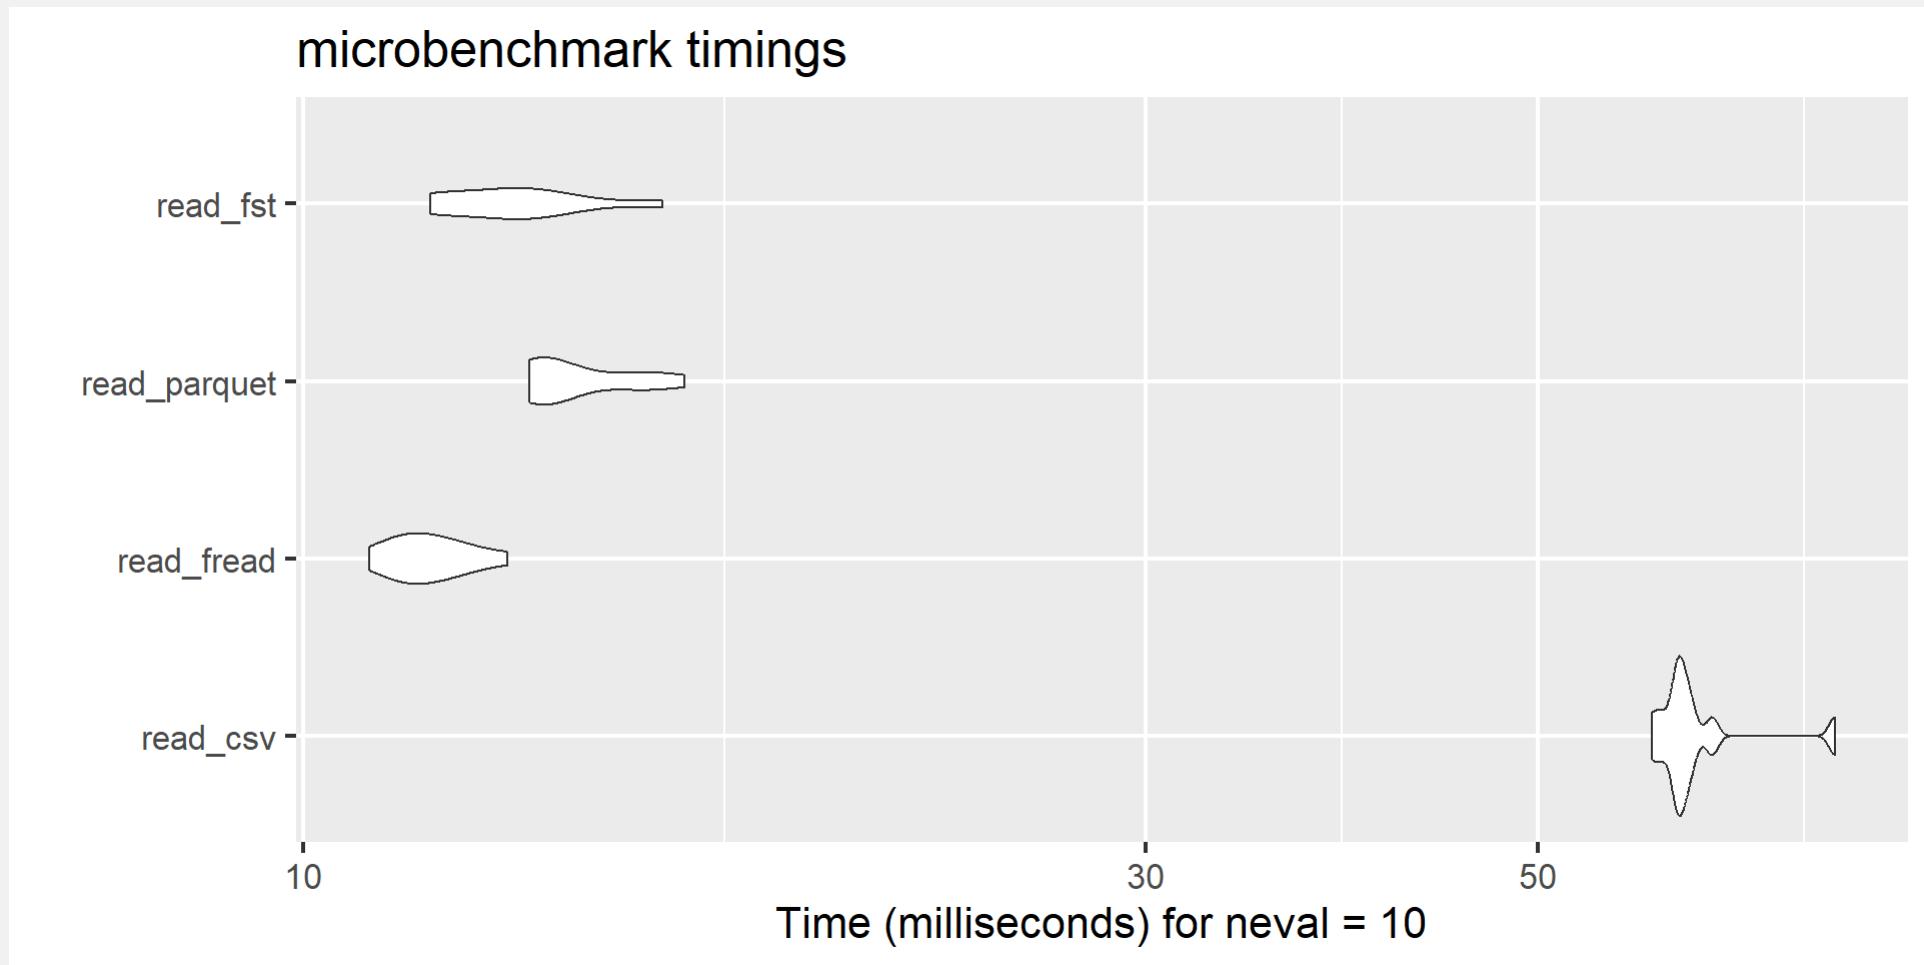
Binary file types can be faster (especially as file size increases)

```
file_path_parquet <- here::here("data/ghg_ems.parquet")
file_path_fst <- here::here("data/ghg_ems.fst")

compare_read_binary <- microbenchmark(
  # tidyverse
  read_csv = readr::read_csv(file_path_csv, show_col_types = FALSE),
  
  # data.table
  read_fread = data.table::fread(file_path_csv, showProgress = FALSE),
  
  # arrow parquet format
  read_parquet = arrow::read_parquet(file_path_parquet),
  
  # fst format (efficient R only format)
  read_fst = fst::read_fst(file_path_fst),
  
  times = 10
)

autoplot(compare_read_binary)
```

Read data - binary



Write data

Every corresponding read function has a write counterpart:

- `write_csv` from the `readr` package (tidyverse)
- `fwrite` from the `data.table` package
- `write_csv_arrow`, `write_parquet` from the `arrow` package
- `write_fst` from the `fst` package

Write data

```
# create an example with 1000000 rows to export
df <- data.frame(x = 1:1000000, y = 1:1000000, z = 1:1000000)

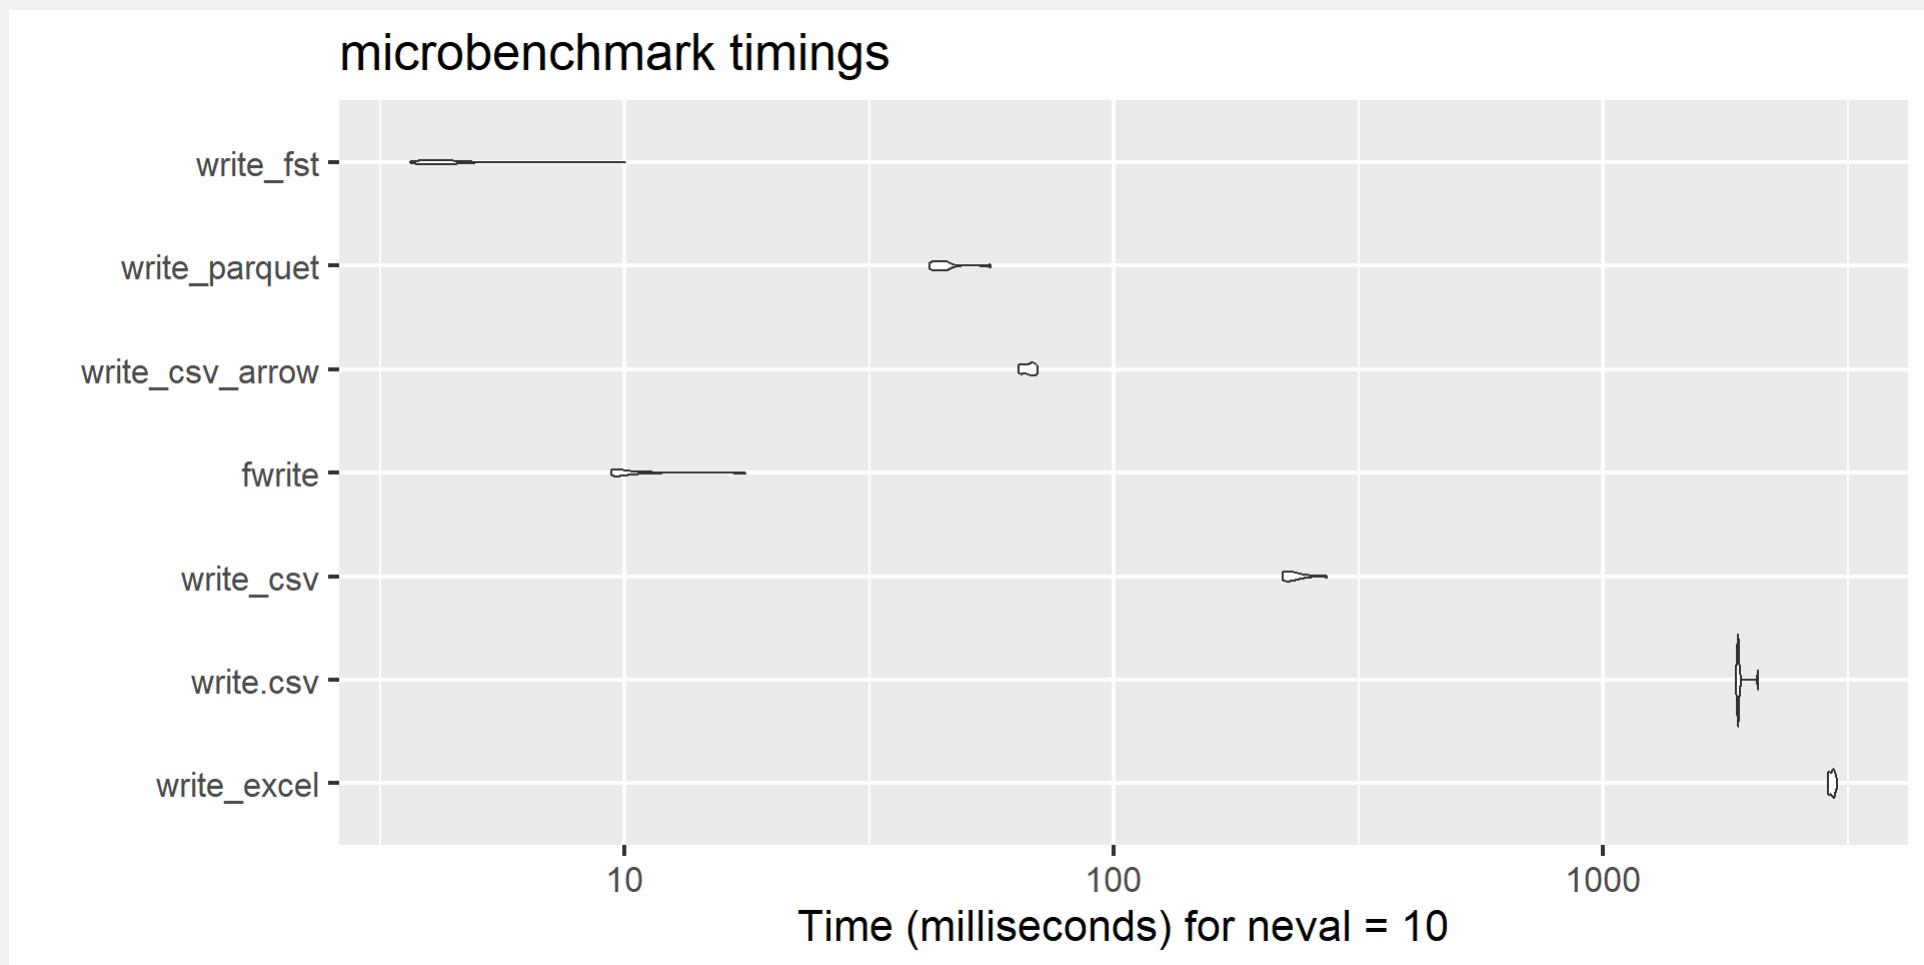
compare_output <- microbenchmark::microbenchmark(
  # write Excel
  write_excel = writexl::write_xlsx(df, "df.xlsx"),

  # write text
  write.csv = write.csv(df, "df.csv"),
  write_csv = readr::write_csv(df, "df.csv"),
  fwrite = data.table::fwrite(df, "df.csv"),
  write_csv_arrow = arrow::write_csv_arrow(df, "df.csv"),

  # write binary formats
  write_parquet = arrow::write_parquet(df, "data/df.parquet"),
  write_fst = fst::write_fst(df, "data/df.fst"),
  times = 10
)

autoplot(compare_output)
```

Write data



Data manipulation

Example: Summarize mean carbon emissions from Electricity by Country

```
library(data.table)
library(dplyr)
library(collapse)
library(arrows)

# Use the right data formats
ghg_emissions <- readr::read_csv(file_path_csv)
ghg_emissions_dt <- setDT(ghg_emissions) # to data table
ghg_emissions_parquet <- read_parquet(file_path_parquet, as_data_frame = FALSE)
```

Summarize data by group

```
1 # 1. The data.table way
2 summarize_dt <- function() {
3   ghg_emissions_dt[, mean(Electricity, na.rm = TRUE), by = Country]
4 }
5
6 # 2. The dplyr way
7 summarize_dplyr <- function() {
8   ghg_emissions |>
9     group_by(Country) |>
10    summarize(mean_e = mean(Electricity, na.rm = TRUE))
11 }
12
13 # 3. The collapse way
14 summarize_collapse <- function() {
15   ghg_emissions |>
16     fgroup_by(Country) |>
17     fsummarise(mean_e = fmean(Electricity))
18 }
19
20 # 4. The arrow way
21 summarize_arrow <- function() {
```

Efficient data manipulation

Example: Summarize mean carbon emissions from Electricity by Country

```
# compare the speed of all versions
compare_summarize <- microbenchmark(
  dplyr = summarize_dplyr(),
  data_table = summarize_dt(),
  collapse = summarize_collapse(),
  arrow = summarize_arrow(),
  times = 10
)

summary(compare_summarize, unit = "relative")
#>      expr      min       lq     mean   median      uq     max neval
#> 1   dplyr  4.792294  4.924788  4.347089  4.578174  4.041363  4.221200    10
#> 2 data_table  2.238343  2.608554  2.469827  2.729258  2.494380  2.175025    10
#> 3   collapse 1.000000  1.000000  1.000000  1.000000  1.000000  1.000000    10
#> 4     arrow 16.002969 16.226567 15.801069 15.930190 14.333536 19.986203    10
#>   cld
#> 1 a
#> 2 ab
#> 3  b
#> 4   c
```

Advanced optimization

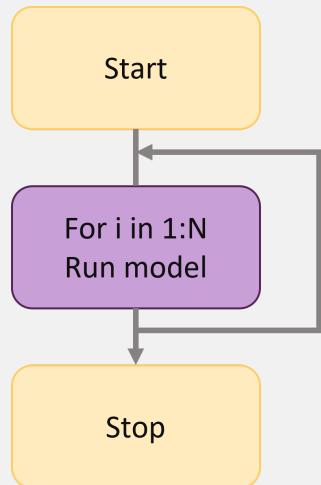
Parallelization and C++

Parallelization

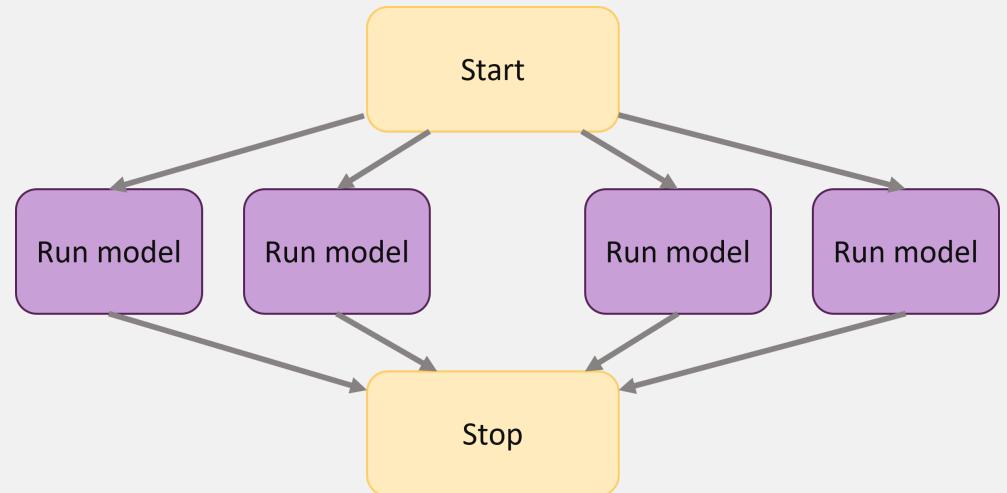
By default, R works on one core but CPUs have multiple cores

```
# Find out how many cores you have with the parallel package  
# install.packages("parallel")  
parallel::detectCores()  
#> [1] 32
```

Sequential



Parallel



Parallelization with the futureverse

- `future` is a framework to help you parallelize existing R code
 - Parallel versions of base R apply family
 - Parallel versions of `purrr` functions
 - Parallel versions of `foreach` loops



A slow example

Let's create a very slow square root function

```
slow_sqrt <- function(x) {  
  Sys.sleep(1) # simulate 1 second of computation time  
  sqrt(x)  
}
```

Before you run anything in parallel, tell R how many cores to use:

```
# Load future package  
library(future)  
# Plan parallel session with 5 cores  
plan(multisession, workers = 5)
```

Parallel apply functions

To run the function on a vector of numbers we could use

Sequential `lapply`

```
# create a vector of 10 numbers
x <- 1:10
tic()
result <- lapply(x, slow_sqrt)
toc()
#> 10.22 sec elapsed
```

Parallel `future_lapply`

```
# Load future.apply package
library(future.apply)

tic()
result <- future_lapply(x, slow_sqrt)
toc()
#> 2.6 sec elapsed
```

Parallel apply functions

Selected base R apply functions and their future versions:

base	future.apply
lapply	future_lapply
sapply	future_sapply
vapply	future_vapply
mapply	future_mapply
tapply	future_tapply
apply	future_apply
Map	future_Map

Parallel for loops

A normal for loop:

```
z <- list()
for (i in 1:10) {
  z[i] <- slow_sqrt(i)
}
```

First, change into a **foreach**

```
library(foreach)
z <- foreach(i = 1:10) %do%
{
  slow_sqrt(i)
}
```

Parallel for loops

Use `doFuture` and `foreach` package to parallelize for loops

The **sequential** version

```
library(foreach)

tic()
z <- foreach(i = 1:10) %do%
{
  slow_sqrt(i)
}
toc()
#> 10.17 sec elapsed
```

The **parallel** version

```
library(doFuture)

tic()
z <- foreach(i = 1:10) %dofuture%
{
  slow_sqrt(i)
}
toc()
#> 3.11 sec elapsed
```

Future purrr functions

The `furrr` package offers parallel versions of `purrr` functions

The **sequential** version

```
library(purrr)

# the purrr version
tic()
z <- map(x, slow_sqrt)
toc()
#> 10.11 sec elapsed
```

The **parallel** version

```
library(furrr)

# the furrr version
tic()
z <- future_map(x, slow_sqrt)
toc()
#> 3.02 sec elapsed
```

Close multisession

When you are done working in parallel, explicitly close your multisession:

```
# close the multisession plan  
plan(sequential)
```

Replace slow code with C++

- Use the [Rcpp package](#) to re-write R functions in C++
- [Rcpp](#) is also used internally by many R packages to make them faster
- Requirements:
 - Install C++ compiler (Rtools for Windows, Xcode for Mac, gcc for Linux, see [here](#) for instructions)
 - Some knowledge of C++
- See [this book chapter](#) and the [online documentation](#) for more info

Rewrite a function in C++

Example: R function to calculate Fibonacci numbers

```
# A function to calculate Fibonacci numbers
fibonacci_r <- function(n) {
  if (n < 2) {
    return(n)
  } else {
    return(fibonacci_r(n - 1) + fibonacci_r(n - 2))
  }
}
```

```
# Calculate the 30th Fibonacci number
fibonacci_r(30)
#> [1] 832040
```

Rewrite a function in C++

Use `cppFunction` to rewrite the function in C++:

```
library(Rcpp)

# Rewrite the fibonacci_r function in C++
fibonacci_cpp <- cppFunction(
  "int fibonacci_cpp(int n){
    if (n < 2){
      return(n);
    } else {
      return(fibonacci_cpp(n - 1) + fibonacci_cpp(n - 2));
    }
  }"
)
```

```
# calculate the 30th Fibonacci number
fibonacci_cpp(30)
#> [1] 832040
```

Rewrite a function in C++

You can also source C++ functions from C++ scripts.

C++ script `fibonacci.cpp`:

```
#include "Rcpp.h"

// [[Rcpp::export]]
int fibonacci_cpp(const int x) {
  if (x < 2) return(x);
  return (fibonacci(x - 1)) + fibonacci(x - 2);
}
```

Then source the function in your R script using `sourceCpp`:

```
sourceCpp("fibonacci.cpp")

# Use the function in your R script like you are used to
fibonacci_cpp(30)
```

How much faster is C++?

```
compare_rcpp <- microbenchmark(
  r = fibonacci_r(30),
  rcpp = fibonacci_cpp(30),
  times = 10
)

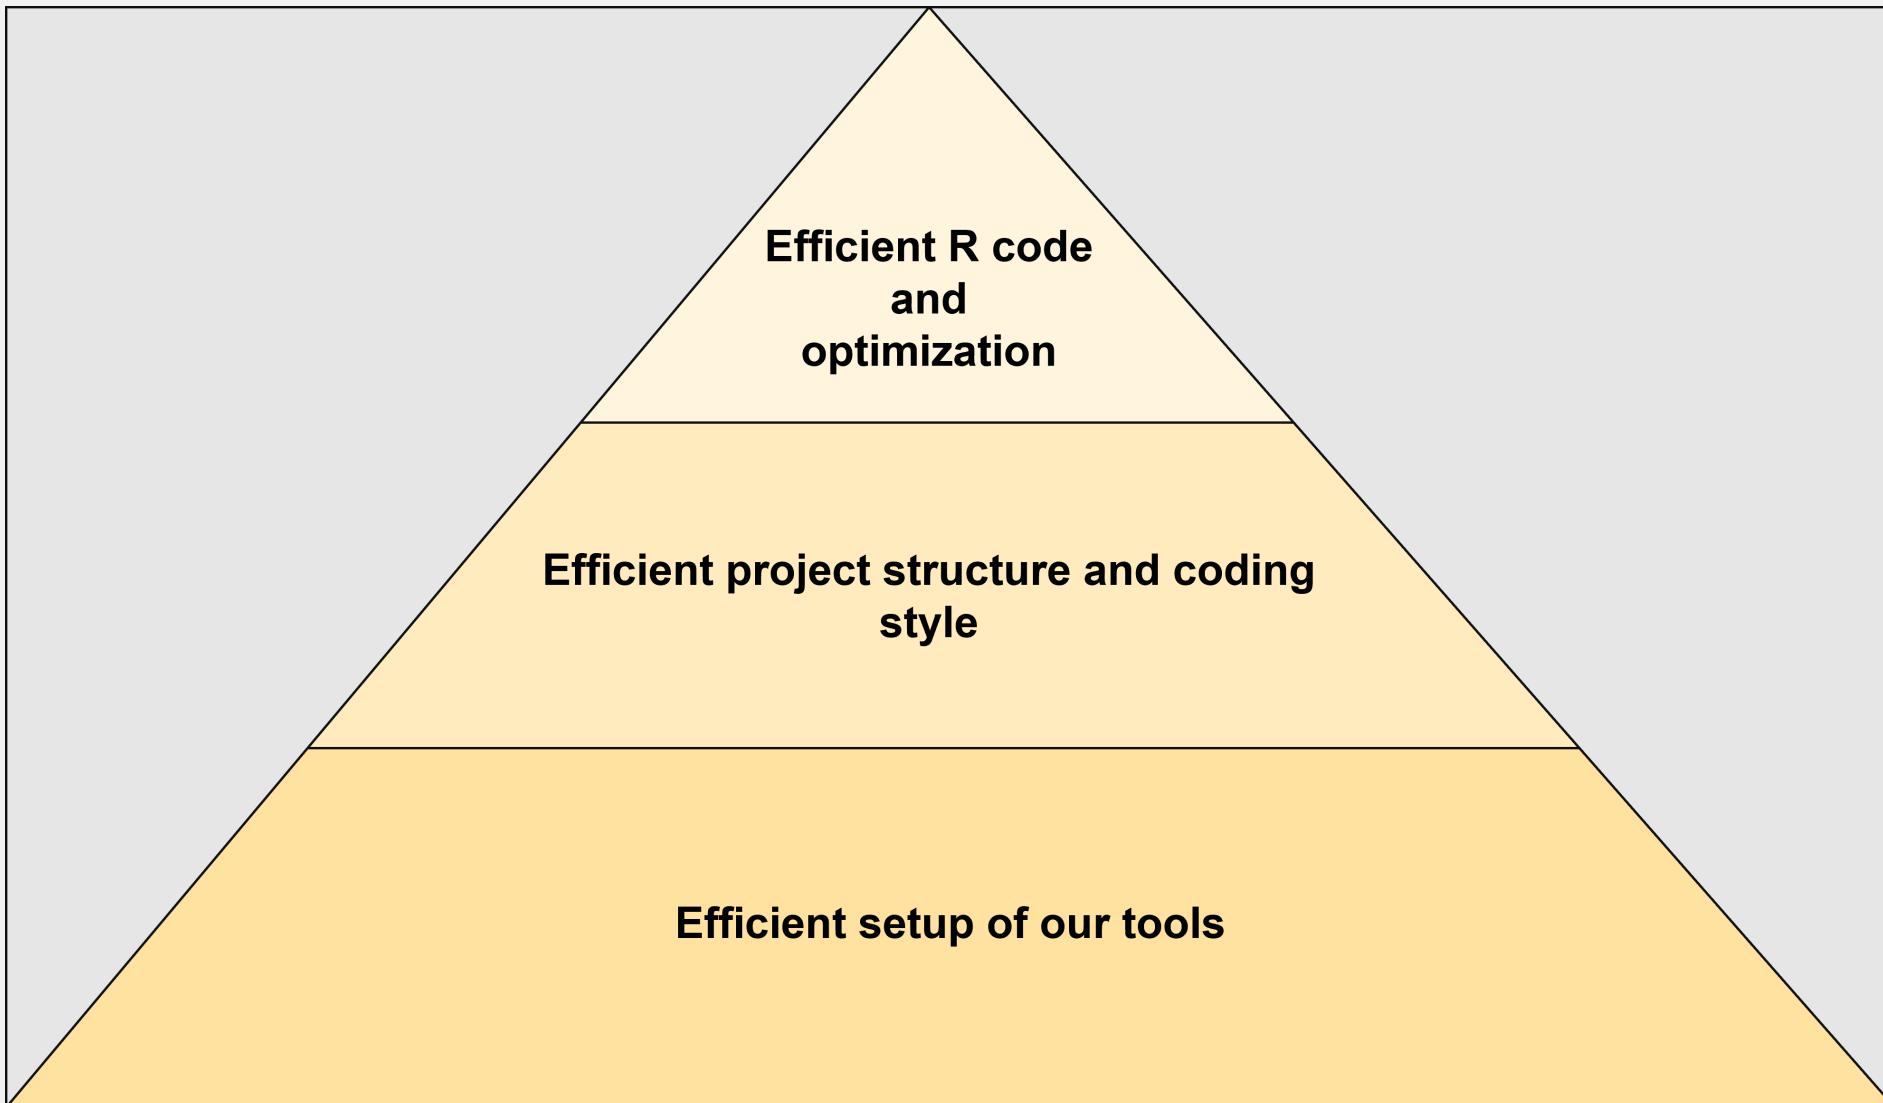
summary(compare_rcpp, unit = "relative")
#>   expr      min      lq     mean    median      uq     max neval cld
#> 1   r 457.9179 462.7536 448.4766 465.1011 467.8475 365.7407     10  a
#> 2 rcpp  1.0000  1.0000   1.0000   1.0000   1.0000   1.0000     10  b
```

Summary

Efficient R code and optimization

- Find your code bottlenecks and optimize those
 - Use `profvis`, `microbenchmark`, and `tictoc`
- Many quick wins, e.g.
 - Use `fread` instead of `read.csv`
 - Use `future_apply/future_map` instead of `apply/map`
 - Replace some tidyverse/base R functions with collapse functions
- If necessary: Advanced optimization techniques (e.g. C++ integration)
- For memory bottlenecks:
 - Consider data bases together with `dbplyr`
 - Use `arrow` to work on larger-than-RAM datasets

Summary



Thank you for your attention

Questions?

Appendix

Cache function results

- Use the `memoise` package
- If functions are called many times with the same arguments
 - Avoids the recalculation
- Useful to e.g. improve the performance of a shiny app

Cache function results

Example: Create a plot on a subset of the `iris` data set

```
# Example of using memoise to cache results
library(memoise)
library(ggplot2)

# Remove rows from plotting function
select_iris_species <- function(rows_to_remove) {
  iris_subset <- iris[-rows_to_remove, ]
  # Do a plot on the subset
  p <- ggplot(
    iris_subset,
    aes(x = Sepal.Length, y = Sepal.Width, color = Species))
  ) +
  geom_point()
}

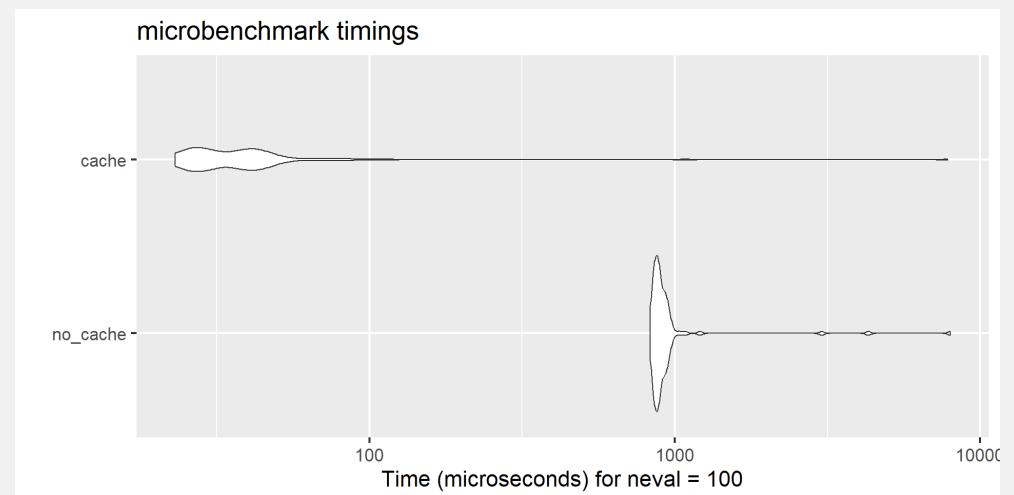
# Version of the function with memoise
select_iris_species_mem <- memoise(select_iris_species)
```

Cache function results

Example: Create a plot on a subset of the `iris` data set

```
# Compare the two versions
result <- microbenchmark(
  no_cache = select_iris_species(10),
  cache = select_iris_species_mem(10)
)

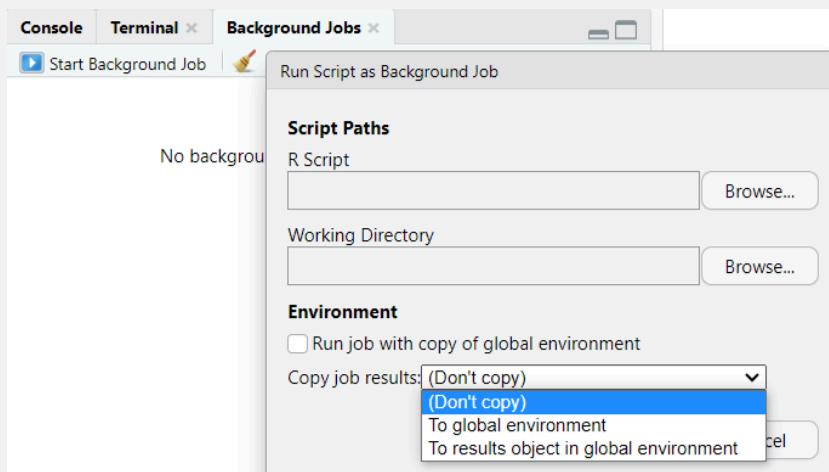
autoplot(result)
```



Run the code somewhere else

- Run it on a cluster
- Run your script in the background

Use RStudio background jobs



Start your R script from the command line

```
Rscript my_script.R
```